

Highlights

Rethinking Decoder Redundancy in Efficient 3D Medical Image Segmentation: A Characterization Study

Author One, Author Two, Author Three

- Decoder contribution is spatially heterogeneous and highly concentrated.
- Efficient decoders achieve net-neutral voxel correctness at matched parameters.
- Uncertainty signals predict where predictions flip, but not improvement direction.
- Decoder redundancy is structured, enabling future adaptive inference strategies.

Rethinking Decoder Redundancy in Efficient 3D Medical Image Segmentation: A Characterization Study

Author One^{a,*}, Author Two^b, Author Three^a

^a*Department Name, University Name, Address Line, City Name, 00000, State, Country*

^b*Another Department, Another University, Address Line, City Name, 00000, State, Country*

Abstract

Recent advancements in U-Net based 3D medical image segmentation have introduced increasingly computation-heavy decoders, particularly with the integration of attention and Transformer-based modules. In response, efficient models such as EffiDec3D simplify or remove these decoders based on empirical end-point tradeoffs and the assumption that static decoder capacity is globally redundant. However, these efficiency-focused approaches fail to answer a fundamental scientific question: where, how much, and under what conditions is this heavy decoder computation actually useful? To fill this gap, we present a systematic characterization study of decoder computation. By evaluating voxel-level disagreements between matched full and efficient decoders across representative architectures and datasets, we characterize decoder utility in terms of heterogeneity, concentration, and predictability. Our results demonstrate that decoder computation is spatially heterogeneous; the full decoder provides a statistically net-neutral benefit, with improvements counterbalanced by degradations. However, this net-neutral effect is highly concentrated in a thin boundary band, revealing structured redundancy rather than uniform irrelevance. Finally, we show that while cheap inference-time uncertainty signals reliably predict the location of prediction instability, they fail to predict the direction of improvement, rendering simple selective-allocation strategies ineffective. Ultimately, our findings reveal

*Corresponding author

Email address: `author1@institute.edu` (Author One)

that modern decoder computation exhibits structured redundancy, providing the scientific constraints necessary for designing future adaptive and dynamic decoding architectures.

Keywords: 3D Medical Image Segmentation, Efficient Deep Learning, Uncertainty, Decoder Redundancy, Model Characterization

1. Introduction

Precise 3D medical image segmentation is critical for clinical applications such as disease screening, surgical planning, and image-guided intervention. To meet these demands, deep learning architectures have historically relied on the encoder-decoder paradigm established by U-Net. Over the past several years, the field has witnessed a continuous evolution toward increasingly expressive encoder representations. Concurrently, modern architectures—including UNETR, Swin UNETR, 3D UX-Net, and MedNeXt—have largely retained dense, multi-resolution decoding stages. More recent designs further enhance decoder capability through advanced feature fusion, refinement, and upsampling modules. This trajectory reflects an implicit assumption that heavy decoder computation is universally beneficial for accurate segmentation.

Despite these continuous advances, existing literature evaluates decoder design almost exclusively through aggregate, end-point metrics such as the Dice similarity coefficient or Hausdorff distance. Such global metrics offer a macroscopic view of model performance and efficiency trade-offs, but they critically obscure the micro-level behavior of the network. Specifically, they fail to answer a fundamental scientific question: *Where, how much, and under what conditions is additional decoder computation actually useful?*

To address this gap, we present a systematic characterization study of decoder computation in 3D medical image segmentation. Rather than introducing a new architecture, we employ matched pairs of full and efficient decoders (specifically building upon the EffiDec3D framework) as controlled experimental instruments. By isolating the decoder capacity while holding the encoder, data, and training protocols constant, we directly observe the voxel-level disagreements—the “flips”—caused by adding decoder capacity.

Our characterization framework yields the following core contributions that challenge the prevailing assumption of uniform decoder utility:

- **Heterogeneous Contribution:** We demonstrate that the marginal benefit of a full decoder over an efficient baseline is almost entirely localized to thin boundary bands and small, difficult anatomical structures.
- **Concentrated Benefit:** Our oracle analysis reveals that the vast majority of the full decoder’s positive corrections are clustered in a small fraction of the image volume ($< 5\%$), providing direct evidence of structured decoder redundancy.
- **Predictability of Instability:** We show that while lightweight inference-time uncertainty signals (e.g., entropy) can reliably predict *where* predictions will be unstable, they fail to predict the *direction* of improvement.
- **Net-Neutral Voxel Correctness:** At the voxel level, the matched full decoder’s net benefit is statistically indistinguishable from zero, representing a net-neutral effect where positive corrections are counterbalanced by negative degradations.

Collectively, these findings demonstrate that modern decoder computation exhibits a structured pattern of heterogeneous contribution, concentrated benefit, and partially predictable instability, providing the scientific foundation for future adaptive and dynamic decoding strategies. The rest of this paper is organized as follows: Section 2 reviews related work, Section 3 details our characterization framework, Section 4 presents the experimental setup and results, Section 5 discusses the implications, and Section 6 concludes the paper.

2. Related Work

2.1. Decoder Design in 3D Medical Image Segmentation

U-Net established the dominant encoder-decoder paradigm for medical image segmentation, characterizing dense, multi-stage decoding and skip connections. Subsequent architectural advancements have consistently evolved toward increasingly expressive decoder computation. Modern networks, including UNETR, Swin UNETR, 3D UX-Net, and MedNeXt, have substantially strengthened encoder representations while largely retaining dense multi-resolution decoder designs. More recent architectures, such as Swin DER,

further enhance decoder capability through improved upsampling, feature fusion, and feature refinement mechanisms.

Despite these continuous advances in decoder design, prior work evaluates decoders almost exclusively through final segmentation accuracy and aggregate computational cost. Little work investigates *where* additional decoder computation contributes spatially, or whether its contribution is uniform across the target anatomy.

2.2. *Efficient Decoder Design*

In response to the computational bottleneck of 3D segmentation, recent works have begun to explicitly target decoder efficiency. The EffiDec3D framework [1] introduced a systematic approach to decoder simplification through channel reduction and the removal of high-resolution decoder stages, establishing a strong architecture-level efficiency-accuracy tradeoff. Similarly, EfficientMedNeXt [2] adopted comparable decoder simplifications before introducing specialized convolution blocks.

Other lightweight segmentation architectures improve efficiency through general architectural redesign, lightweight attention mechanisms, or hybrid encoders (e.g., UNETR++, Slim UNETR, SegFormer3D) rather than explicitly reducing or characterizing decoder computation. Crucially, existing efficient decoder studies infer redundancy implicitly from end-point accuracy-efficiency trade-offs. They do not directly characterize where decoder computation contributes at the voxel level.

2.3. *Characterization Studies*

Existing characterization studies in medical image segmentation primarily investigate prediction uncertainty, calibration, prediction reliability, failure detection, and segmentation quality prediction [3, 4, 5, 6]. These studies often evaluate how well uncertainty metrics correlate with overall error rates or how they can be used to flag uncertain cases for clinician review.

Our work fundamentally diverges from this literature. Rather than characterizing general prediction uncertainty, we characterize the marginal contribution of decoder computation itself. Prior work has reduced decoder computation or studied prediction uncertainty, but has not systematically characterized the utility of decoder capacity through voxel-level heterogeneity, benefit concentration, and inference-time predictability.

3. Methods: A Characterization Framework

To understand the spatial and functional utility of decoder computation, we design a controlled observation study. We treat a matched full decoder and an efficient decoder as experimental instruments, isolating the effect of decoder capacity on voxel-level predictions.

3.1. Experimental Setup

Let v denote a voxel within a 3D medical image. We evaluate two decoder configurations sharing identical encoder features, training protocols, and data splits:

- **Efficient Decoder (E_1)**: Parameter-matched efficient decoder based on the EfficDec3D configuration, utilizing concatenation for skip aggregation to strictly align with published parameter counts.
- **Full Decoder (E_0)**: The standard dense, high-resolution decoder of the respective backbone architecture.

For a given voxel v , let $y(v)$ be the ground truth label, $\hat{y}_e(v)$ be the prediction from the efficient decoder, and $\hat{y}_f(v)$ be the prediction from the full decoder.

3.2. Voxel-level Flip Analysis

Rather than relying on aggregate metrics like Dice, which sum all errors indiscriminately, we track how the full decoder alters the prediction of the efficient baseline. We define three categories of prediction “flips” when transitioning from the efficient to the full decoder:

- **Positive Flip ($P(v)$)**: The full decoder improves a previously incorrect prediction. $P(v) = 1[\hat{y}_e(v) \neq y(v) \wedge \hat{y}_f(v) = y(v)]$.
- **Negative Flip ($N(v)$)**: The full decoder degrades a previously correct prediction. $N(v) = 1[\hat{y}_e(v) = y(v) \wedge \hat{y}_f(v) \neq y(v)]$.
- **Net Benefit ($U(v)$)**: The voxel-wise difference between improvements and degradations. $U(v) = P(v) - N(v)$.

The global rates are defined as $R_{pos} = \text{mean}_v P(v)$, $R_{neg} = \text{mean}_v N(v)$, and $R_{net} = R_{pos} - R_{neg}$.

3.3. Core Metrics

We operationalize our investigation into structured redundancy through the following specific metrics:

1. Global Net Benefit: This metric evaluates the macro-level necessity of the full decoder. The global net flip rate R_{net} is computed as the expected voxel-wise net benefit over the volume V :

$$R_{net} = \frac{1}{|V|} \sum_{v \in V} U(v) \quad (1)$$

Its physical meaning answers whether the full decoder provides a unidirectional advantage ($R_{net} > 0$) or if its positive corrections and negative degradations cancel each other out ($R_{net} \approx 0$).

2. Boundary and Anatomy Resolution: To evaluate spatial heterogeneity, we analyze the net benefit $U(v)$ as a function of the Euclidean distance to the nearest organ boundary, as well as its distribution across different anatomical structures. This reveals *where* the decoder is actually intervening.

3. Oracle Concentration: This metric measures the spatial sparsity of the decoder’s utility. We divide the image into 3D blocks b_r and compute the net benefit $B(r) = \sum_{v \in b_r} U(v)$ for each block. We rank the blocks by $B(r)$ and calculate the fraction of the total positive net mass recovered by the top $k\%$ of blocks:

$$Coverage(k) = \frac{\sum_{r=1}^{k\% \times |B|} \max(B(r), 0)}{\sum_{r=1}^{|B|} \max(B(r), 0)} \quad (2)$$

We summarize this curve using the K_{80} metric, which is the execution budget required to recover 80% of the full decoder’s positive benefit.

4. Flip Predictability: This evaluates whether we can forecast the decoder’s behavior before executing it. We use lightweight inference-time signals from the efficient decoder, such as predictive entropy:

$$H(v) = - \sum_{c=1}^C p_c(v) \log p_c(v) \quad (3)$$

where $p_c(v)$ is the softmax probability for class c . We measure predictability using the Area Under the Receiver Operating Characteristic (AUROC) to

predict the *location* of flips (instability) and the *direction* of flips (positive vs. negative).

5. Selection Gap: This measures the practical, deployable opportunity for selective computation. We compare the net benefit recovered by an uncertainty-based block selector (e.g., selecting blocks with the highest mean entropy) against a random selection policy at a fixed computational budget (e.g., 20%). A significant positive gap indicates a viable routing strategy.

4. Experiments and Results

4.1. Experimental Setup

We evaluate matched decoder pairs across representative datasets to answer three core questions: (1) Is the full decoder’s contribution spatially uniform or heterogeneous? (2) Is the benefit broadly distributed or concentrated in specific regions? (3) Can we predict where the full decoder will improve predictions using cheap inference-time signals?

To ensure cross-architecture generality, our characterization study evaluates a large-kernel CNN (3D UX-Net) and a transformer-based architecture (SwinUNETR). Experiments are conducted on the BTCV multi-organ CT dataset (13 organs) and the FeTA 2021 fetal brain MRI dataset (7 structures).

Hardware and Environment: All models were trained on an AutoDL RTX 5090 (32 GB VRAM, Blackwell architecture) utilizing CUDA 12.8 and PyTorch 2.6. We employed BF16 mixed precision to accelerate training.

Training Protocol: Both the full decoder models (E_0) and the efficient baseline models (E_1) were trained for a total of 45,000 optimizer steps with identical data augmentations, learning rates (0.001), and a crop size of 96×96 . Model evaluation was performed every 250 steps.

Reproduction and Controlled Settings: To serve as a rigorous controlled baseline, our efficient decoder implementation strictly utilizes canonical concatenation for skip aggregation, matching the exact published parameter count of EffiDec3D (e.g., 3.15M parameters for 3D UX-Net E_1 , compared to 53.0M for the full E_0). On the BTCV dataset, our E_1 baseline achieves a mean Dice of 77.73%. While this is slightly below the original paper’s reported 79.25% (primarily due to variations in the preprocessing pipeline and dataset splits), the tight control over the encoder weights, data split, and training regimen between our E_0 and E_1 pairs perfectly isolates the true spatial effect of the decoder computation for our observation study.

4.2. Heterogeneity: The Net-Neutral Boundary Band

We first evaluate the overall voxel-level net benefit of the full decoder. Surprisingly, with the parameter-matched efficient decoder configuration, we find that the full decoder’s net benefit is statistically indistinguishable from zero. For 3D UX-Net on BTCV, the subject-mean net flip rate is +0.046% with a 95% confidence interval spanning $[-0.0002\%, +0.0011\%]$. This indicates a “net-neutral” voxel correctness: the full decoder corrects some predictions but degrades an almost equal number of previously correct predictions.

However, this net-neutral behavior is highly heterogeneous. The boundary error rate is $3.93\times$ higher than the interior error rate. When mapping the net flip rate as a function of distance to the organ boundary (Figure 2), the residual flip activity is sharply concentrated in a thin boundary band (0-1 voxels from the edge).

Furthermore, decoder activity is anatomically concentrated. Small and morphologically complex structures experience significantly higher flip rates than large organs. For example, the Left and Right Adrenal glands exhibit positive flip rates of 16.2% and 16.9% respectively, but equally severe negative flip rates, yielding highly turbulent, net-negative impacts (e.g., -7.1% net benefit for the L.Adrenal). Conversely, large organs like the liver see minimal flip activity ($< 3\%$). This demonstrates that while the full decoder provides negligible net voxel benefit globally, it remains influential specifically at critical organ boundaries and within difficult structures.

4.3. Concentration: Structured Redundancy

We investigate the spatial concentration of decoder benefit using an oracle block-selector. Figure 3 presents the oracle coverage curve. The curve is extremely steep: an oracle selecting only the top 5% of blocks recovers 80% of the full decoder’s positive net mass ($K_{80} = 5\%$). By a 10% budget, the oracle recovers 99.98%, and near 100% at 20%. This stark concentration provides direct evidence of structured decoder redundancy, proving that the vast majority of the image volume derives no meaningful benefit from the heavy decoder layers.

4.4. Predictability: Instability vs. Improvement

Given the concentrated nature of decoder benefit, we evaluated whether lightweight inference-time signals (specifically, predictive entropy and confidence from the efficient decoder) could identify these regions prior to executing the full decoder.

Our results reveal a stark dichotomy between predicting instability and predicting improvement. The entropy signal is highly predictive of flip *locations* (instability), achieving a subject-mean AUROC of 0.903 (95% CI: [0.876, 0.926]). However, when tasked with predicting the *direction* of the flip (positive vs. negative), the AUROC collapses to 0.592 (95% CI: [0.560, 0.626]), performing only marginally better than random chance.

Consequently, the deployable selection gap is zero. When we use entropy to allocate a 20% computation budget to the most “uncertain” blocks, the net benefit recovered is 0.288, which is statistically indistinguishable from a random allocation policy (0.059). The 95% CI for the paired difference crosses zero ($[-0.059, 0.451]$). Confidence behaves identically. This proves that simple uncertainty thresholding cannot reliably recover the theoretical efficiency opportunity.

5. Discussion

The results of our characterization study compel a re-evaluation of how decoder computation is allocated and assessed in 3D medical image segmentation. The central finding—that the marginal benefit of a full decoder is net-neutral at the voxel level and tightly clustered around organ boundaries—challenges the necessity of dense, uniform decoding for the entirety of an input volume.

Crucially, our study delineates the boundary between observation and action. The high AUROC for flip location confirms that the regions where predictions change are deterministic and easily identified by cheap uncertainty signals. We observe that high entropy strongly correlates with boundary regions, exactly where the decoder acts. However, the inability to predict the *direction* of these flips (improvement versus degradation) represents a significant negative result for simple region-routing strategies. Because the baseline efficient decoder is already highly capable, deploying the full decoder on uncertain regions is just as likely to disturb a correct prediction as it is to fix an incorrect one.

This implies that we cannot realize the theoretical computational savings of structured redundancy merely by selectively gating the execution of existing static decoders using uncertainty thresholds. The predictability of instability (but not improvement) necessitates a shift from attempting to “turn off” decoder layers to developing genuinely adaptive architectures that

dynamically allocate capacity based on feature complexity rather than simple confidence scores.

6. Conclusion

In this work, we presented a systematic characterization study of decoder computation in modern 3D medical image segmentation architectures. By analyzing the voxel-level differences between matched full and efficient decoders, we demonstrated that modern decoder computation exhibits structured redundancy: its contribution is highly heterogeneous, its benefits are spatially concentrated at object boundaries, and the instability of its predictions is easily detectable.

However, we also showed that the directional utility of this computation is currently unpredictable using standard inference-time signals. The full decoder is net-neutral in voxel correctness, revealing that simple uncertainty-based routing is insufficient to harness this redundancy. This work establishes the inherent limitations of static decoding and naive thresholding. Future work will leverage these constraints to design genuinely dynamic architectures that adaptively allocate decoder computation based on local feature complexity, moving beyond simple confidence scores to harness this structured redundancy.

CRediT authorship contribution statement

Author One: Conceptualization, Methodology, Writing – original draft. **Author Two:** Data curation, Software. **Author Three:** Supervision, Writing – review & editing.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Data availability

The data used in this study (BTCV/Synapse, FeTA 2021) are publicly available.

References

- [1] A. Author, et al., Efficient 3d medical image segmentation, Journal Name (2024).
- [2] B. Author, et al., Efficientmednext: Multi-receptive dilated convolutions for medical image segmentation, Journal Name (2024).
- [3] A. Jungo, M. Reyes, Assessing reliability and challenges of uncertainty estimations for medical image segmentation, MICCAI (2019).
- [4] C. Author, et al., Evaluating uncertainty quantification in medical image segmentation: A multi-dataset, multi-algorithm study, Computerized Medical Imaging and Graphics (2025).
- [5] D. Author, et al., Uncertainty-aware segmentation quality prediction via deep learning bayesian modeling, Computerized Medical Imaging and Graphics (2025).
- [6] E. Author, et al., A review of uncertainty estimation and its application in medical imaging, Journal Name (2023).




Figure 1: Motivation: Decoder computation yields spatially localized flips. While the full decoder corrects some efficient decoder errors (positive flips), it simultaneously degrades other previously correct predictions (negative flips) at high-entropy boundaries. This illustrates that decoder computation is not uniformly beneficial.




Figure 2: Heterogeneity. Left: Global positive and negative flip rates demonstrating net-neutral behavior. Middle: Net flip rate distributed by distance to the boundary. Right: Organ-level flip rates showing high volatility in small structures (e.g., adrenal glands).




Figure 3: Concentration. Oracle coverage curve showing that 80% of the full decoder's net benefit is concentrated in just 5% of the image volume.




Figure 4: Predictability. Entropy highly discriminates flip locations (instability, AUROC 0.903) but fails to distinguish positive corrections from negative degradations (AUROC 0.592).